AI Evals for Designers: Quick Checklist

A five-step check for designers. About two hours with 20 real answers from your product, and nothing beyond a spreadsheet.

01 **Collect 20 real answers**

Pull 20 real conversations or outputs from your product. Include the awkward ones: off-topic requests, missing data, things it should refuse.

In the study: The study used 41 screenshots and deliberately included pricing, login and legal pages, where an AI audit has nothing to find.

02 **Mark each one pass or fail**

Ask one question: would I be happy if a real user got this? Write one plain line on the first thing wrong. No scores and no categories yet.

In the study: Each of the 41 audits got one verdict and, when it failed, one line naming the first thing wrong.

03 **Group the notes and count**

Sort the notes into piles the way you would affinity-map research. Count each pile. The biggest one is your first problem.

In the study: Six piles. Two tied for biggest at 7 failed audits each: stating guesses as findings, and asking for confidence scores on ordinary chat.

04 **Write the biggest problem down as one rule**

Make it a yes or no question that anyone on the team could apply to an answer and get the same verdict.

In the study: The rule: every piece of advice is either seen (it points at something on the screen) or assumed. 93 of 116 pieces of advice were assumed.

05 **Fix, then re-check the same 20 plus 5 new ones**

Change one thing, rerun the same examples and apply the same rule. Add a few answers the fix never saw, so you know it works beyond the ones you tuned it on.

In the study: Guessed advice fell from 80% to 68% on the tuned screens, but only from 92% to 88% on fresh ones.

#	Input (what the user asked or did)	Output (what the AI said)	Pass / Fail	First thing wrong (one line)	Group
1					
2					
3					
4					
5					
6					
7					
8					
9					
10					

Full method: aiuxdesign.guide/guides/ai-evals-for-designers

Review sheet: aiuxdesign.guide/ai-evals-checklist